

Introduction to Hugging Face



Juan David Martínez Vargas
jdmartinev@eafit.edu.co

2023

What is Hugging Face?



- 1.NLP & Machine Learning Company:** Leading innovator in the field of natural language processing.
- 2.Transformers Library:** Offers a user-friendly interface to state-of-the-art models like BERT, GPT-2, and T5.
- 3.Model Hub:** A platform for sharing and discovering pre-trained NLP models.
- 4.Datasets Library:** Centralized access to diverse NLP datasets.
- 5.Community-Driven:** Active forums, collaborations, and open-source contributions.

Bridging the gap between research and real-world applications in NLP

What is Hugging Face?

The screenshot shows the GitHub repository for Hugging Face's `transformers` library. The repository is public and has 14,050 commits, 221 branches, and 136 tags. The latest commit by `natolambert` is titled "Add torch RMSProp optimizer (#26425)". The repository includes a file tree with folders like `.circleci`, `.github`, `docker`, `docs`, `examples`, `model_cards`, `notebooks`, and `scripts`. The `scripts` folder contains a file `transformers-cli` that links to `huggingface-cli`. The right sidebar provides an "About" section describing the library as "State-of-the-art Machine Learning for Pytorch, TensorFlow, and JAX" and lists various supported models and frameworks like `python`, `nlp`, `machine-learning`, `tensorflow`, `pytorch`, `transformer`, `speech-recognition`, `seq2seq`, `flax`, `pretrained-models`, `language-models`, `nlp-library`, `language-model`, `hacktoberfest`, `bert`, `jax`, `pytorch-transformers`, and `model-hub`.

huggingface / transformers

Code Issues 627 Pull requests 205 Actions Projects 25 Security Insights

transformers Public

Watch 1.1k Fork 22.3k Star 112k

main 221 branches 136 tags

Go to file Add file <> Code

natolambert Add torch RMSProp optimizer (#26425) 408b2b3 2 hours ago 14,050 commits

File	Commit	Time
.circleci	Add Nougat (#25942)	15 hours ago
.github	Integrate AMD GPU in CI/CD environment (#26007)	last week
docker	Integrate AMD GPU in CI/CD environment (#26007)	last week
docs	Deleted duplicate sentence (#26394)	12 hours ago
examples	feat: adding num_proc to load_dataset (#26326)	4 days ago
model_cards	Update URL for Hub PR docs (#17532)	last year
notebooks	Update README.md (#25941)	3 weeks ago
scripts	transformers-cli -> huggingface-cli (#23276)	4 months ago

About

Transformers: State-of-the-art Machine Learning for Pytorch, TensorFlow, and JAX.

huggingface.co/transformers


python nlp machine-learning natural-language-processing deep-learning tensorflow pytorch transformer speech-recognition seq2seq flax pretrained-models language-models nlp-library language-model hacktoberfest bert jax pytorch-transformers model-hub

What is Natural Language Processing

1. **Definition:** NLP stands for Natural Language Processing.
2. **Bridge Between Humans & Computers:** Enables machines to understand, interpret, and generate human language.
3. **Applications:**
 1. Text and speech recognition
 2. Machine translation
 3. Sentiment analysis
 4. Chatbots and virtual assistants
4. **Combines:** Linguistics, computer science, and artificial intelligence to simulate human language abilities in machines.

Making machines understand and respond to us more naturally

Pipeline




```
from transformers import pipeline

classifier = pipeline("sentiment-analysis")
classifier("I've been waiting for a HuggingFace course my whole life.")
```

```
[{'label': 'POSITIVE', 'score': 0.9598047137260437}]
```

The pipeline function returns an end-to-end object that performs an NLP task on one or several texts

seq2seq models



```
{'id': ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'],  
  'translation': [{'es': '¡Intentemos algo!',  
                    'pt': 'Vamos tentar alguma coisa!'},  
                  {'es': '¡Intentemos algo!', 'pt': 'Vamos tentar algo!'},  
                  {'es': 'Tengo que irme a dormir.', 'pt': 'Preciso ir dormir.'},  
                  {'es': 'Tengo que irme a dormir.', 'pt': 'Tenho que ir dormir.'},  
                  {'es': 'Tengo que irme a dormir.', 'pt': 'Tenho de dormir.'},  
                  {'es': '¿Qué estás haciendo?', 'pt': 'O que está fazendo?'},  
                  {'es': '¿Qué estás haciendo?', 'pt': 'O que você está fazendo?'},  
                  {'es': '¿Qué estás haciendo?', 'pt': 'O que estás a fazer?'},  
                  {'es': '¿Qué es eso?', 'pt': 'O que é aquilo?'},  
                  {'es': '¿Qué es eso?', 'pt': 'O que é isso?'}]}
```

Available pipelines

- feature-extraction (get the vector representation of a text)
- fill-mask
- ner (named entity recognition)
- question-answering
- sentiment-analysis
- summarization
- text-generation
- translation
- zero-shot-classification

```
from transformers import pipeline
```

```
classifier = pipeline("zero-shot-classification")  
classifier(  
    "This is a course about the Transformers library",  
    candidate_labels=["education", "politics", "business"],  
)
```

```
{  
    'sequence': 'This is a course about the Transformers library',  
    'labels': ['education', 'business', 'politics'],  
    'scores': [0.8445963859558105, 0.111976258456707, 0.043427448719739914]}  
}
```

Pre-processing - tokenizers


Word-based

Character-based

Subword-based

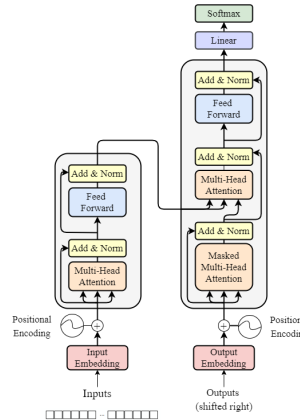
Pre-processing - tokenizers

Raw text
This is an
amazing talk



Neural networks can't process
raw text 

Convert text into numbers 



1. Split text into words, subwords or symbols (punctuation) – tokens
2. Map each token to an integer
3. Add additional inputs

Pre-processing - tokenizers

Word-based

Each word has a token (integer) that will turn into a word embedding

A lot of embeddings!!!

Split on spaces

Let's

do

tokenization!

Split on punctuation

Let

's

do

tokenization

!

the	→	1
of	→	2
and	→	3
to	→	4
in	→	5
was	→	6
the	→	7
is	→	8
for	→	9
as	→	10
on	→	11
with	→	12
that	→	13
dog	→	14
dogs	→	15

Pre-processing - tokenizers

Character-based

Each character has a token (integer) that will turn into a word embedding

Less embeddings 😊

Less meaningful 😞

L	e	t	'	s	d	o	t	o	k	e	n	i	z	a	t	i	o	n	!
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Pre-processing - tokenizers

Subword-based

Frequently used words are should not be split into smaller subwords

Rare words should be decomposed into meaningful subwords



Let's </w>

do</w>

token

ization</w>

!</w>



Vocabulary

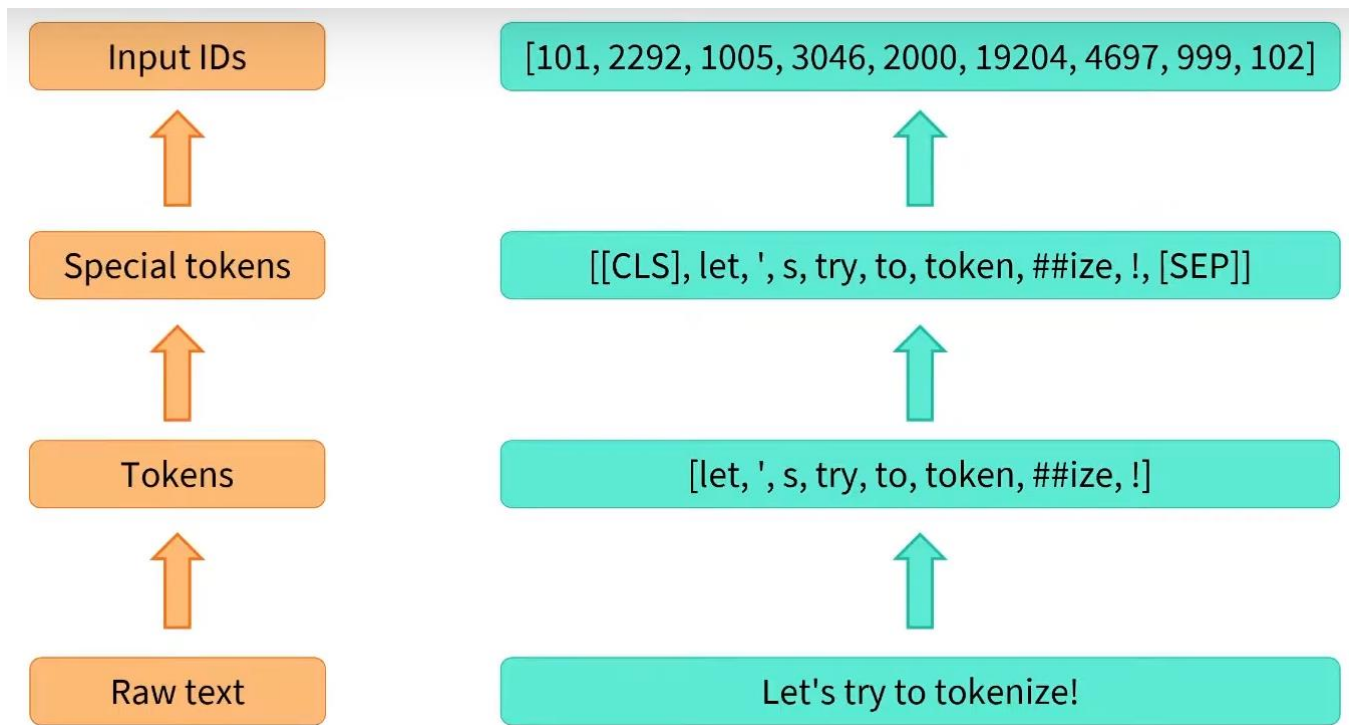
Text	Label
The ghost pepper is so spicy, it is hauntingly hot	1
I tried to hug the sun today, but it was too hot to handle	1
I cannot handle spicy food	0



Vocabulary

but
cannot
food
ghost
handle
hauntingly
hot
hug
i
is
it
pepper
so
spicy
sun
the
to
today
too
tried
was

Tokenization pipeline



Hugging Face Tokenizer

```
from transformers import AutoTokenizer
model_checkpoint='t5-small'
tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)

sequence = "Using a Transformer network is simple"
tokens = tokenizer.tokenize(sequence)

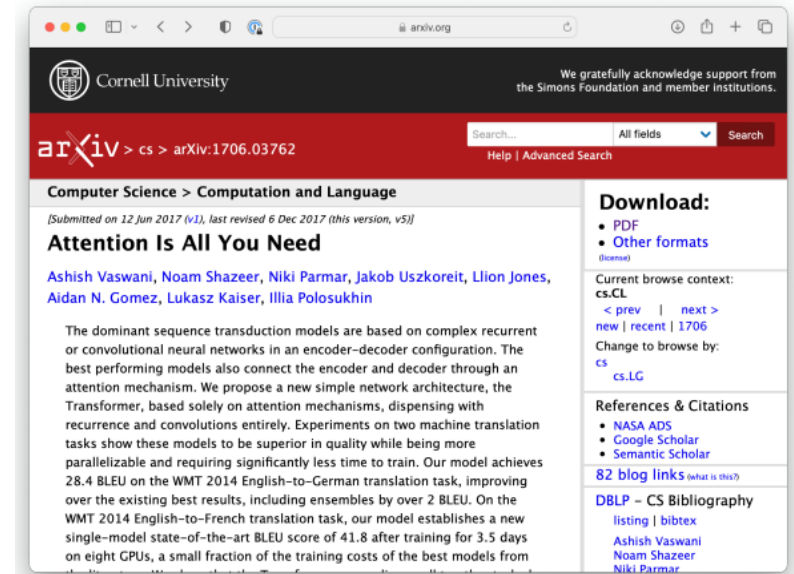
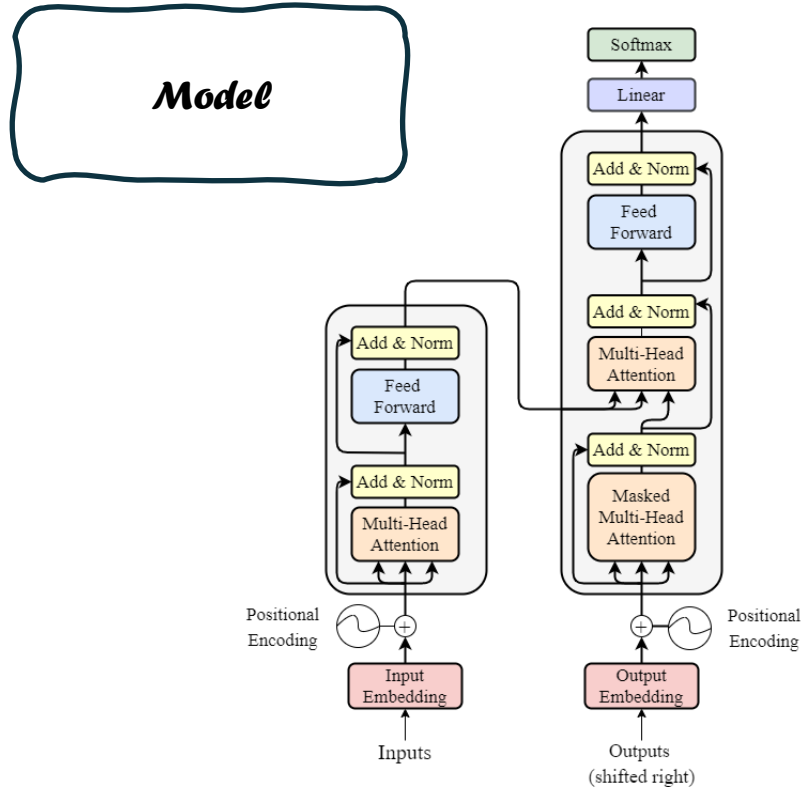
print(tokens)

['_', 'Using', '_', 'a', '__Transformer', '__network', '__is', '__simple']

ids = tokenizer.convert_tokens_to_ids(tokens)
print(ids)


[3, 3626, 3, 9, 31220, 1229, 19, 650]
```

Model - Transformer



Embeddings

```
[[0.6912, 0.8765, 0.4939],  
 [0.6342, 0.7481, 0.7717],  
 [0.8395, 0.2128, 0.3696],  
 [0.4900, 0.1509, 0.0689],  
 [0.2587, 0.9171, 0.8670],  
 [0.7213, 0.9922, 0.5701],  
 [0.7598, 0.5231, 0.3666],  
 [0.5150, 0.5216, 0.9682],  
 [0.2248, 0.0261, 0.4427],
```



```
[101, 2292, 1005, 3046, 2000, 19204, 4697, 999, 102]
```

```
import torch
```

```
torch.manual_seed(123);
```

```
idx = torch.tensor([2, 3, 1]) # 3 training examples
```

```
num_idx = max(idx)+1
```

```
out_dim = 5
```


```
embedding = torch.nn.Embedding(num_idx, out_dim)
```

```
embedding(idx)
```


```
tensor([[ 0.6957, -1.8061, -1.1589,  0.3255, -0.6315],  
        [-2.8400, -0.7849, -1.4096, -0.4076,  0.7953],  
        [ 1.3010,  1.2753, -0.2010, -0.1606, -0.4015]],  
        grad_fn=<EmbeddingBackward0>)
```

Each training example has
5 feature values


Positional encoding




Positional encoding:



Positional vectors:





Input vectors:




Encoder – Attention - Training

Idea: Create a context vector that encodes the meaning of the Word within the text









$$z^{(i)} = \sum_{j=1}^T \alpha_{ij} \cdot x^{(j)}$$




$T = 5$

Decoder – Masked Attention - Training

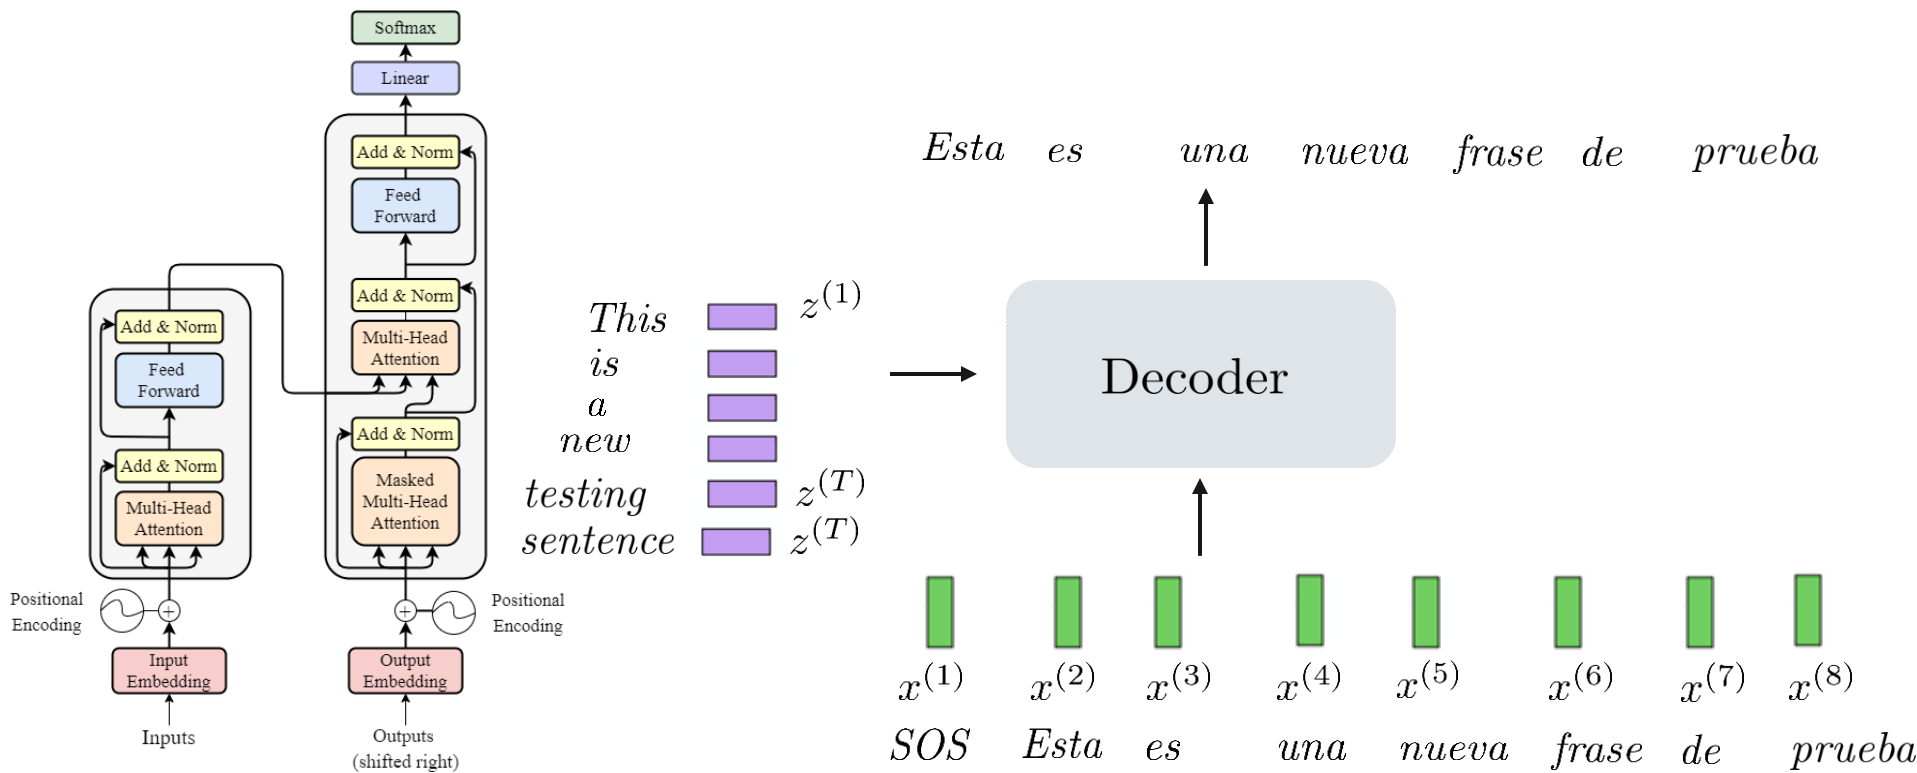


This  $z^{(1)}$
is 
an 
example 
sentence  $z^{(T)}$

Esta es una frase ejemplo
 $p(x^{(1)})$ $p(x^{(2)})$ $p(x^{(3)})$ $p(x^{(4)})$ $p(x^{(5)})$
Esta es la frase ejemplo



Transformer - Deploy



Transformer – Hugging Face

```
from transformers import AutoModelForSeq2SeqLM  
model_checkpoint='t5-small'  
model = AutoModelForSeq2SeqLM.from_pretrained(model_checkpoint)
```

Transfer Learning

Pre-training



Supervised and self supervised learning

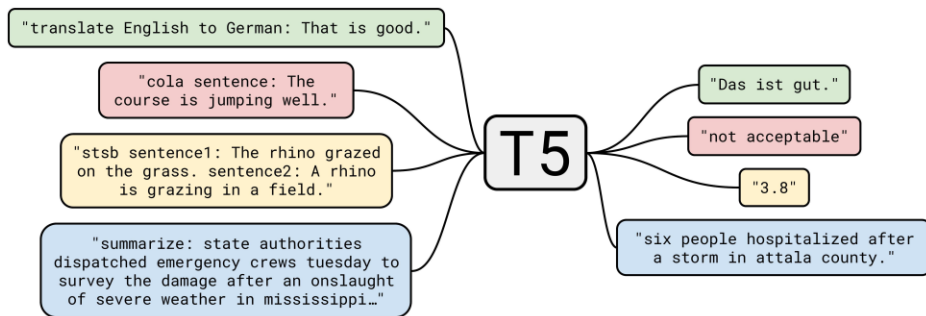


Fine-tuning

```
{'id': ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'],  
'translation': [{'es': 'Intentemos algo!',  
  'pt': 'Vamos tentar alguma coisa!'},  
{'es': 'Intentemos algo!', 'pt': 'Vamos tentar algo!'},  
{'es': 'Tengo que irme a dormir.', 'pt': 'Preciso ir dormir.'},  
{'es': 'Tengo que irme a dormir.', 'pt': 'Tenho que ir dormir.'},  
{'es': 'Tengo que irme a dormir.', 'pt': 'Tenho de dormir.'},  
{'es': '¿Qué estás haciendo?', 'pt': 'O que está fazendo?'},  
{'es': '¿Qué estás haciendo?', 'pt': 'O que você está fazendo?'},  
{'es': '¿Qué estás haciendo?', 'pt': 'O que estás a fazer?'},  
{'es': '¿Qué es eso?', 'pt': 'O que é aquilo?'},  
{'es': '¿Qué es eso?', 'pt': 'O que é isso?'}]}
```



Transfer Learning – Transformer T5



Computer Science > Machine Learning

[Submitted on 23 Oct 2019 (v1), last revised 19 Sep 2023 (this version, v4)]

Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, Peter J. Liu

Transfer learning, where a model is first pre-trained on a data-rich task before being fine-tuned on a downstream task, has emerged as a powerful technique in natural language processing (NLP). The effectiveness of transfer learning has given rise to a diversity of approaches, methodology, and practice. In this paper, we explore the landscape of transfer learning techniques for NLP by introducing a unified framework that converts all text-based language problems into a text-to-text format. Our systematic study compares pre-training objectives, architectures, unlabeled data sets, transfer approaches, and other factors on dozens of language understanding tasks. By combining the insights from our exploration with scale and our new “Colossal Clean Crawled Corpus”, we achieve state-of-the-art results on many benchmarks covering summarization, question answering, text classification, and more. To facilitate future work on transfer learning for NLP, we release our data set, pre-trained models, and code.

Questions

Thank you

